LMU LUDWIG-MAXIMILIANS-UNIVERSITÄT MÜNCHEN

A Hidden Gem: Unlocking the Power of Bayesian Inference in Particle Physics

Lorenz Gärtner[1], Toni Sculac, Judith Katzy

[1]*LMU Munich*

February 17, 2025

TUM LMU ORIGINS Excellence Cluster

Belle II

PUNCH 4NFDI

DFG

About me

- BSc @ University of Manchester
  Physics with theoretical physics
  Little stats
- MSc @ LMU Munich
  Thesis on QFT in curved spacetime
  Almost no stats
- Currenly PhD @ LMU Munich
  A lot of stats

What is a probability?

Kolmogorov probability axioms

1. $p(\Omega) = 1$, where Ω is the sample space.
2. $p(x) \geq 0$ for any event $x \subseteq \Omega$.
3. For any sequence of disjoint events $x_1, x_2, \ldots,$

$$p\left(\bigcup_{i=1} x_i\right) = \sum_{i=1} p(x_i)$$

Conditional probability

Is **defined** as the probability of an event x if we know that an event y is true $p(x|y)$.

$$p(x \cup y) = p(x|y)p(y) \quad \rightarrow \quad p(x|y) = \frac{p(x \cup y)}{p(y)}$$

Probability interpretations

The probability axioms and rules allow you to calculate new probabilities from old ones.

To assign probabilities we need probability interpretations.

The interpretations **share the same mathematical framework**, but the meaning of $p(x)$ is different.

Frequentist interpretation

Assign a probability as relative frequency

$$p(x) = \lim_{n \to \infty} \frac{n_x}{n}$$

- Can only be assigned to repeatable experiments
- Not everything is repeatable...
- No probabilities of single events

Bayesian interpretation

Assign a probability $p(x)$ as *degree of belief*.

- Probability depends on the experimenters' knowledge.
- Inference results are subjective.
- Everything is a random variable.

Bayes' theorem

The **posterior** is

$$p(\text{theory}|\text{data}) = \frac{p(\text{data}|\text{theory})p(\text{theory})}{p(\text{data})}$$

- **Likelihood** $p(\text{data}|\text{theory})$
- **Prior** $p(\text{theory})$
- **Evidence** $p(\text{data}) = \int p(\text{data}|\text{theory})p(\text{theory})$

Can you derive it?

The common ground

Frequentist inference is based on

- the model of the observable data

$$\boldsymbol{p(x|\theta)}$$

Bayesian inference is based on

- the model of the observable data

$$\boldsymbol{p(x|\theta)}$$

- your prior belief

$$p(\theta)$$

➜ You want the best possible $\boldsymbol{p(x|\theta)}$.

The best possible model...

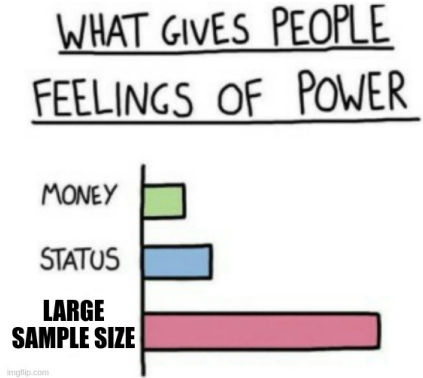

Physics does not care...

... about our interpretation

For many inference problems, the frequentist and Bayesian approaches give similar numerical values, even though they answer different questions.

BUT if results are different, you should understand why.

Parameter inference

Parameter inference

Point estimates

Identify the most probable parameter point.

Interval estimation

Identify extended regions in parameter space based on compatibility with the data.

Frequentist point estimates: estimators

Estimator is a *statistic* $\hat{\theta}(x)$, i.e. a well chosen function of the data with the desited properties of

- *consistency*
  $\lim_{n_x \to \infty} E_x[\hat{\theta}] = \theta_{true}$
- *unbiasedness*
  $b = E_x[\hat{\theta}] - \theta_{true}$
- *efficiency*

$$V(\hat{\theta}) = I(\theta)^{-1} = E_x\left[\left(\frac{\partial \ln p(x|\theta)}{\partial \theta}\right)^2\right]^{-1}$$

- ...

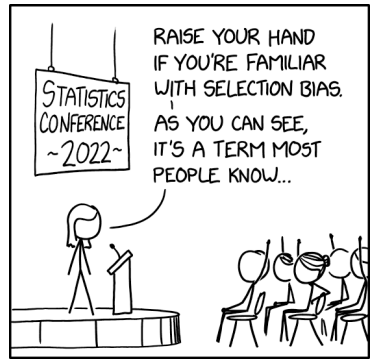

A good data set is the key to success ...

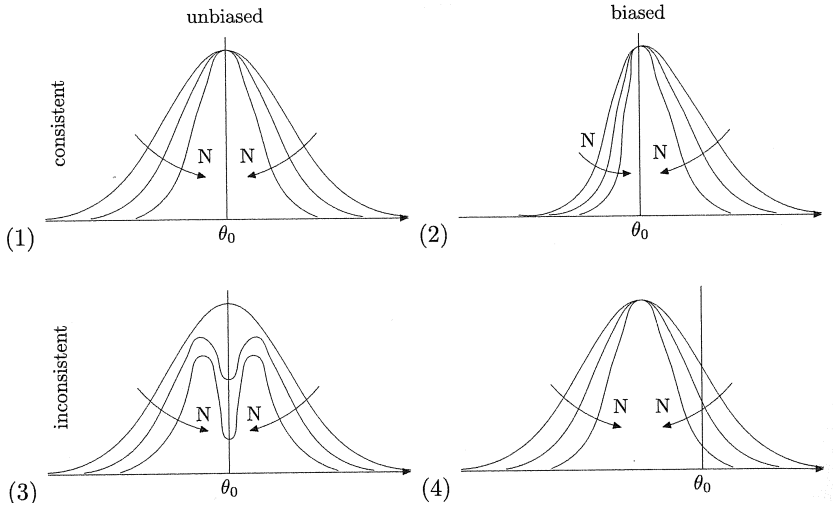

Method of maximum likelihood

We find maximum likelihood estimators $\hat{\theta}$ by solving

$$\left.\frac{\partial p(x|\theta)}{\partial \theta}\right|_{\theta=\hat{\theta}} = 0$$

With the property

$$\lim_{n\to\infty} p\left(\sqrt{n}(\hat{\theta} - \theta_{true})\right) = \mathcal{N}\left(0, I(\theta)^{-1/2}\right)$$

This implies consistency, asymptotic unbiasedness and efficiency.

Bayesian point estimates

Mode
Value of θ with maximum a-posteriori probability

$$\theta^* = \text{argmax}_\theta \, p(\theta|x)$$

Mean
Expected value of θ under the posterior

$$\bar{\theta} = E_{p(\theta|x)}[\theta]$$

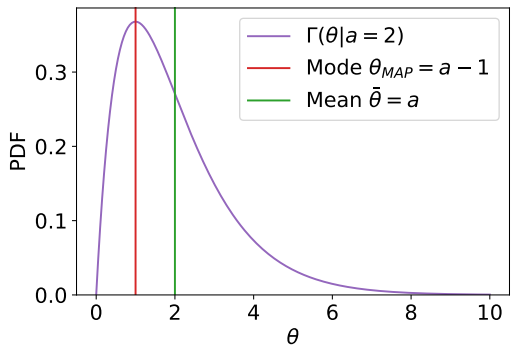

Intervals and limits

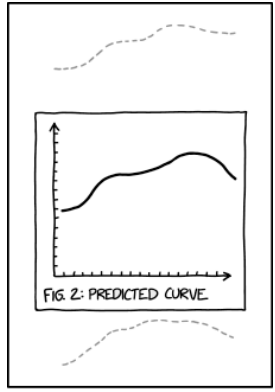

SCIENCE TIP: IF YOUR MODEL IS BAD ENOUGH, THE CONFIDENCE INTERVALS WILL FALL OUTSIDE THE PRINTABLE AREA.

Intervals and limits

If an estimator PDF is not Normal, $\hat{\theta} \pm \sigma_\theta$ is meaningless.

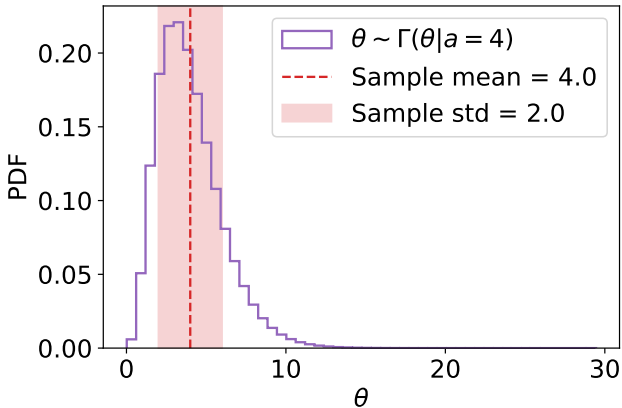

Frequentist intervals

Neyman confidence belt

$$\int_{x_1}^{x_2} dx\ p(x|\theta) = 1 - \alpha$$

Not unique ➜ *central interval*

$$\int_{-\infty}^{x_1} dx\ p(x|\theta) = \int_{x_2}^{\infty} dx\ p(x|\theta) = \alpha/2$$

or *upper/lower interval*

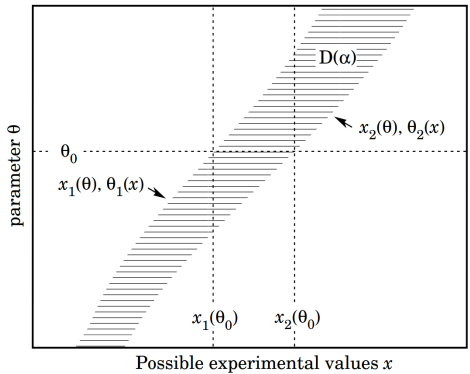

Bayesian intervals

Credible intervals $[\theta_1, \theta_2]$ cover $1 - \alpha$ of the posterior

$$\int_{\theta_1}^{\theta_2} d\theta \; p(\theta|x) = 1 - \alpha$$

- For upper/lower limits: set θ_1 or θ_2 to boundary
- *Smallest possible interval*
  ➜ Highest (posterior) density intervals (HDI)

Nuisance parameters and priors

Frequentist nuisance parameters

Models are not perfect ➜ **systematic bias**
Solution: Include additional **nuisance** parameters ν,

$$p(x|\theta, \nu).$$

Usually, we want to constrian nuisance parameters, but in the frequentist language **everything is data**.

Frequentist "priors"

Constrain nuisance parameters using *auxiliary data a*,

$$p(x|\theta,\nu)p(a|\nu)$$

$p(a|\nu)$ represents our *degree of belief* in ν.

Often *auxiliary data* is *created* to match our desired constraint term.

Bayesian nuisance parameters

Priors can be defined using auxiliary data

$$p(\nu|a) \propto p(a|\nu)p_0(\nu)$$

Only in the Bayesian case, other prior choices are also allowed, e.g.

$$p(\nu) = \text{Gauss}(\nu|\nu_0, \sigma_\nu)$$

Bayesian updating

Can generally use measurements to *update* our posterior

$$p(\theta|x_1, x_2) = \frac{p(x_2|\theta)}{p(x_2)} p(\theta|x_1) = \frac{p(x_2|\theta)}{p(x_2)} \frac{p(x_1|\theta)}{p(x_1)} p(\theta).$$

Effectively what is done for combining measurement results (eg. PDG averages).

Bayesian updating

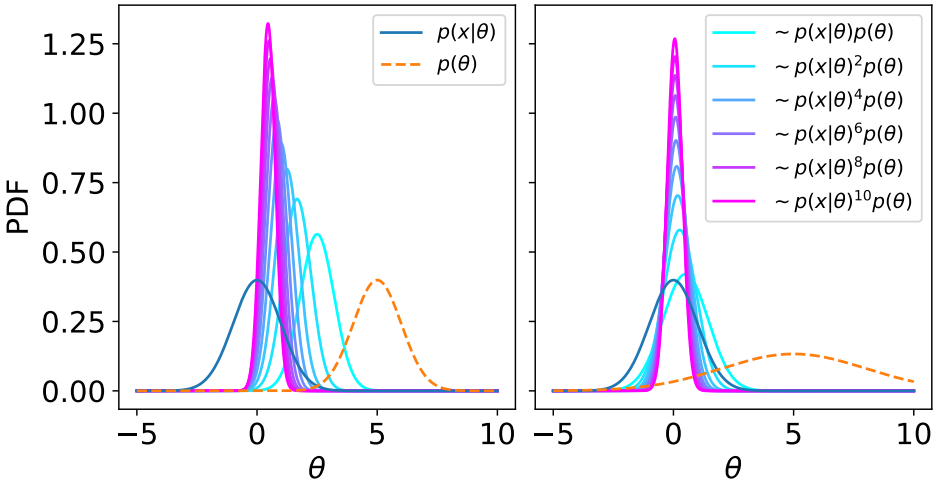

A simple linear model

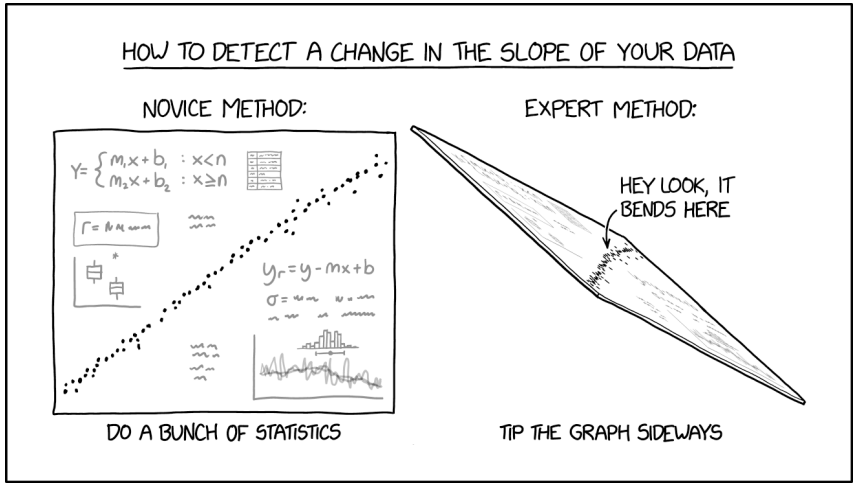

A simple linear model

- Our independent data : $\boldsymbol{X} = (x_i, y_i, \sigma_i)$

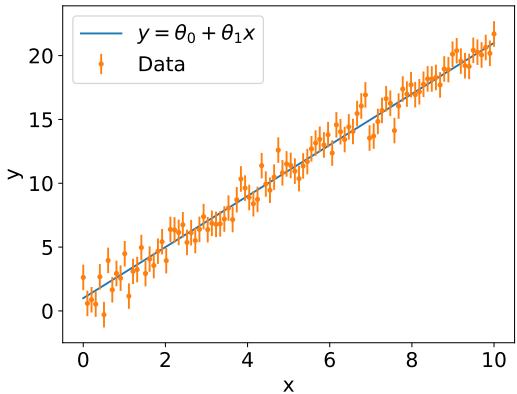

A simple linear model

- Our independent data: $\boldsymbol{X} = (x_i, y_i, \sigma_i)$
- Our model:

$$p(\boldsymbol{X}|\theta_0, \theta_1) = \prod_{x_i, y_i, \sigma_i \in \boldsymbol{X}} \text{Gauss}(y_i|\mu(x_i|\theta_0, \theta_1), \sigma_i)$$

$$\mu(x_i|\theta_0, \theta_1) = \theta_0 + \theta_1 x_i$$

➜ We want to know about θ_0, do not care about θ_1.

Unconstrained likelihood

$$-2 \log p(\boldsymbol{X}|\theta_0, \theta_1) = \sum_{x_i, y_i, \sigma_i \in \boldsymbol{X}} \frac{(y_i - \mu(x_i|\theta_0, \theta_1))^2}{\sigma_i^2}$$

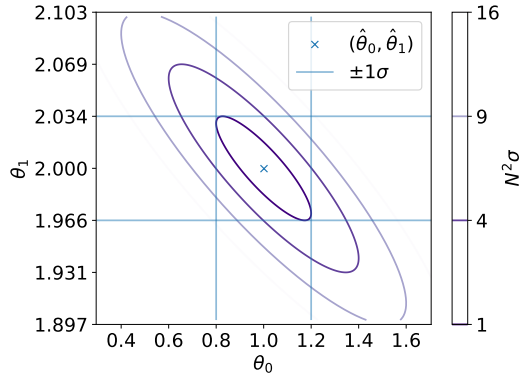

Including a measurement of θ_1: t_1, σ_{t_1}

$$-2 \log p(\boldsymbol{X}|\theta_0, \theta_1) = \sum_{x_i, y_i, \sigma_i \in \boldsymbol{X}} \frac{(y_i - \mu(x_i|\theta_0, \theta_1))^2}{\sigma_i^2} + \frac{(\theta_1 - t_1)^2}{\sigma_{t_1}^2}$$

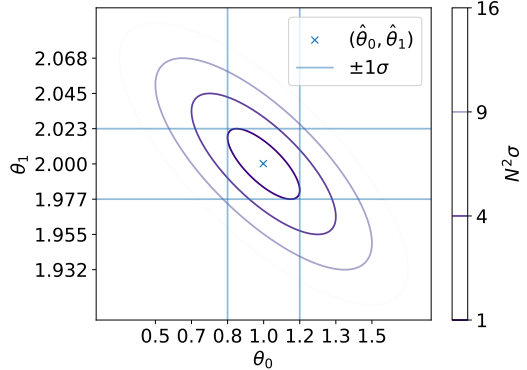

Posterior

$$p(\theta_0, \theta_1 | \boldsymbol{X}) \propto p(\boldsymbol{X} | \theta_0, \theta_1) p(\theta_0) p(\theta_1)$$

$$p(\theta_0) = \text{Uniform}(0, 2)$$

$$p(\theta_1) = \text{Gauss}(\theta_1 | t_1, \sigma_{t_1})$$

Corner plots are great for visualization.

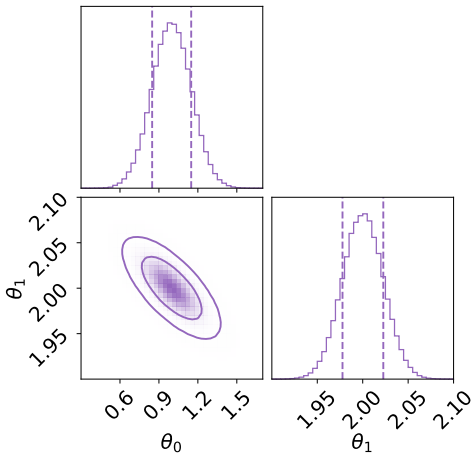

Marginal posterior

$$p(\theta_0|\boldsymbol{X}) = \int d\theta_1 \, p(\theta_0, \theta_1|\boldsymbol{X}) = \text{Gauss}(\theta_0|\theta_0^*, \sigma_{\theta_0})$$

In this example, we get

- $\theta_0^* = \hat{\theta}_0$
- 68% HDI = $\hat{\theta}_0 \pm \sigma_{\theta_0}$

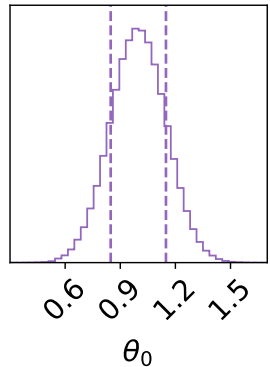

MCMC

The hard part ...

- We only want the posterior for θ alone.
- Remove nuisance parameters by integrating over ν.

➜ The *marginal posterior* is

$$p(\theta|x) = \int d\nu \; p(\theta, \nu|x)$$

- Commonly a high dimensional integral
  ➜ compute with Monte Carlo methods.

Markov chain

A sequence of events, where probability of the next state depends solely on the current state

$$\ldots \to \theta_i \sim g(\theta_i|\theta_{i-1}) \to \theta_{i+1} \sim g(\theta_{i+1}|\theta_i) \to \ldots$$

for some *proposal distribution* g.

Markov Chain Monte Carlo (MCMC)

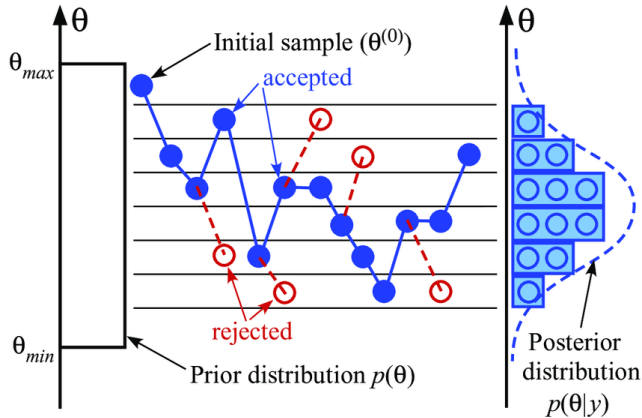

Metropolis-Hastings

We loop

1. Generate $\theta \sim g(\theta|\theta_i)$
2. Update

$$\theta_{i+1} = \begin{cases} \theta & u \leq \min\left(1, \frac{p(\theta)g(\theta|\theta_i)}{p(\theta_i)g(\theta_i|\theta)}\right) \\ \theta_i & \text{otherwise} \end{cases}$$

where $u \sim \text{Uniform}(0, 1)$

Note: Need to define a proposal distribution $g(\theta|\theta_0)$.

Chains

In MCMC we generate a sequence

$$\theta_0 \to \theta_1 \to \theta_2 \to \dots$$

Only one start can land you in local minima.

$$\theta_0^0 \to \theta_1^0 \to \theta_2^0 \to \dots$$
$$\theta_0^1 \to \theta_1^1 \to \theta_2^1 \to \dots$$
$$\theta_0^2 \to \theta_1^2 \to \theta_2^2 \to \dots$$
$$\dots$$

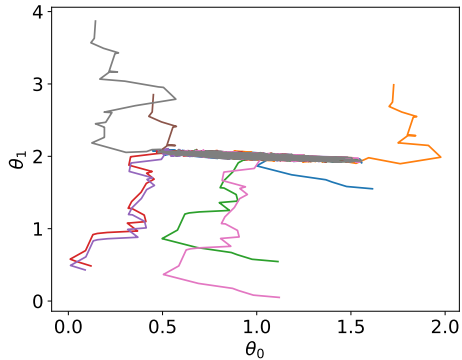

Convergence

Trace plots are a useful convergence diagnostic

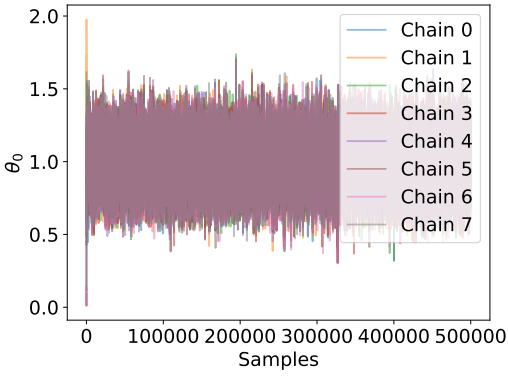

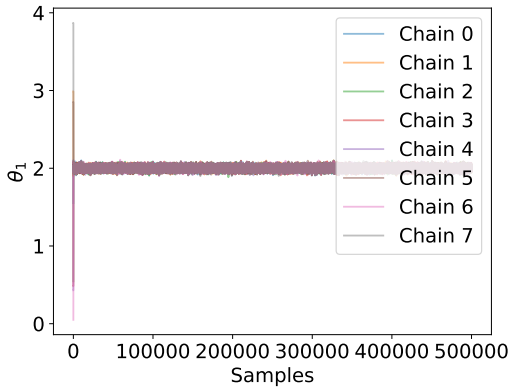

... but one can become more fancy.

Tools to try